

Performance

```
> mongoimport -d school -c students < students.json
```

```
connected to: 127.0.0.1
```

```
2014-05-13T07:12:09.293+0200 check 9 200
```

```
2014-05-13T07:12:09.294+0200 imported 200 objects
```

```
> use school
```

```
switched to db school
```

```
> db.students.findOne()
```

```
{
  "_id" : 0,
  "name" : "aimee Zank",
  "scores" : [
    {
      "type" : "exam",
      "score" : 1.463179736705023
    },
    {
      "type" : "quiz",
      "score" : 11.78273309957772
    },
    {
      "type" : "homework",
      "score" : 6.676176060654615
    },
    {
      "type" : "homework",
      "score" : 35.8740349954354
    }
  ]
}
```

`db.students.find({name:"aimee Zank"})` : parcourt l'intégralité de la collection

Tester avec

```
db.students.find({name:"aimee Zank"}).explain("executionStats")
```

```
"executionStats" : {
  "executionSuccess" : true,
  "nReturned" : 2,
  "executionTimeMillis" : 0,
  "totalKeysExamined" : 0,
  "totalDocsExamined" : 200,
  ...
}
```

```
...
"totalDocsExamined" : 200,
"executionStages" : {
  "stage" : "COLLSCAN",
  "filter" : {
    ...
  }
}
```

`db.students.findOne({name:"aimee Zank"})` : trouve le premier qui correspond, donc plus rapide

Création d'index

```
> db.students.ensureIndex({name:1})
{
  "createdCollectionAutomatically" : false,
  "numIndexesBefore" : 1,
  "numIndexesAfter" : 2,
  "ok" : 1
}
> db.students.find({name:"aimee Zank"}) : plus rapide
```

Tester avec :

`db.students.find({name:"aimee Zank"}).explain("executionStats")`

```
{
  "executionStats" : {
    "executionSuccess" : true,
    "nReturned" : 2,
    "executionTimeMillis" : 0,
    "totalKeysExamined" : 2,
    "totalDocsExamined" : 2,
  }
}
```

...

```
{
  "inputStage" : {
    "stage" : "IXSCAN",
    "nReturned" : 2,
    "executionTimeMillisEstimate" : 0,
  }
}
```

Lister les index d'une collection

```
> db.students.getIndexes()
[
  {
    "v" : 1,
    "key" : {
      "_id" : 1
    },
    "name" : "_id_",
    "ns" : "school.students"
  },
  {
    "v" : 1,
    "key" : {
      "name" : 1
    },
    "name" : "name_1",
    "ns" : "school.students"
  }
]
```

```
    "name" : "name_1",
    "ns" : "school.students"
  }
]
```

Supprimer un index

```
> db.students.dropIndex({name:1})
{ "nIndexesWas" : 2, "ok" : 1 }
```

MultyKey index

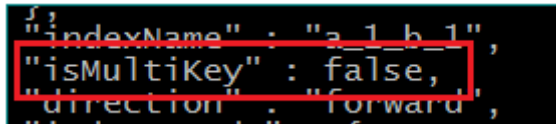
```
mongoimport -d school -c foo < foo.json
```

```
> db.foo.insert({a:1,b:1})
WriteResult({ "nInserted" : 1 })
```

```
> db.foo.ensureIndex({a:1,b:1})
{
  "createdCollectionAutomatically" : false,
  "numIndexesBefore" : 1,
  "numIndexesAfter" : 2,
  "ok" : 1
}
```

```
> db.foo.find({a:1})
{ "_id" : ObjectId("5372490be54ca6f43c383250"), "a" : 1, "b" : 1 }
```

```
> db.foo.find({a:1}).explain("executionStats")
```

A terminal screenshot showing the output of the explain command. The output is a JSON object. The field "isMultiKey" is highlighted with a red rectangle and its value is "false".

```
{
  "indexName" : "a_1_b_1",
  "isMultiKey" : false,
  "direction" : "forward",
  "..."
}
```

```
> db.foo.insert({a:[1,2,3],b:5})
WriteResult({ "nInserted" : 1 })
```

```
> db.foo.find({a:1})
{ "_id" : ObjectId("5372490be54ca6f43c383250"), "a" : 1, "b" : 1 }
{ "_id" : ObjectId("53724b1de54ca6f43c383251"), "a" : [ 1, 2, 3 ], "b" : 5 }
```

```
db.foo.find({a:1}).explain("executionStats")
```

```
{
  "indexName" : "a_1_b_1",
  "isMultiKey" : true,
  "direction" : "forward",
  "indexBounds" : {
```

```
> db.foo.insert({a:[1,2,3],b:[3,4,5]})
WriteResult({
  "nInserted" : 0,
  "writeError" : {
    "code" : 10088,
    "errmsg" : "insertDocument :: caused by :: 10088 cannot index parallel arrays [b] [a]" }
})
```

Option sur la création d'index : unique

```
> db.stuff.insert({thing:"pear"})
WriteResult({ "nInserted" : 1 })

> db.stuff.ensureIndex({thing:1})
{
  "createdCollectionAutomatically" : false,
  "numIndexesBefore" : 1,
  "numIndexesAfter" : 2,
  "ok" : 1
}

> db.stuff.insert({thing:"pear"})
WriteResult({ "nInserted" : 1 })

> db.stuff.insert({thing:"apple"})
WriteResult({ "nInserted" : 1 })

> db.stuff.getIndexes()
[
  {
    "v" : 1,
    "key" : {
      "_id" : 1
    },
    "name" : "_id_",
    "ns" : "test.stuff"
  },
  {
    "v" : 1,
    "key" : {
      "thing" : 1
    },
    "name" : "thing_1",
    "ns" : "test.stuff"
  }
]
```

```
> db.stuff.dropIndex({thing:1})
{ "nIndexesWas" : 2, "ok" : 1 }
```

```
> db.stuff.ensureIndex({thing:1},{unique:1})
{
  "createdCollectionAutomatically" : false,
  "numIndexesBefore" : 1,
  "ok" : 0,
  "errmsg" : "E11000 duplicate key error index: test.stuff.$thing_1  dup key: { : \"pear\" }",
  "code" : 11000
}
```

```
> db.stuff.find()
{ "_id" : ObjectId("53724e24e54ca6f43c383254"), "thing" : "pear" }
{ "_id" : ObjectId("53724e71e54ca6f43c383255"), "thing" : "pear" }
{ "_id" : ObjectId("53724e86e54ca6f43c383256"), "thing" : "apple" }
```

Enlever une collection contenant le mot “**pear**”

```
> db.stuff.ensureIndex({thing:1},{unique:1})
{
  "createdCollectionAutomatically" : false,
  "numIndexesBefore" : 1,
  "numIndexesAfter" : 2,
  "ok" : 1
}
```

```
> db.stuff.getIndexes()
[
  {
    "v" : 1,
    "key" : {
      "_id" : 1
    },
    "name" : "_id_",
    "ns" : "test.stuff"
  },
  {
    "v" : 1,
    "unique" : true,
    "key" : {
      "thing" : 1
    },
    "name" : "thing_1",
    "ns" : "test.stuff"
  }
]
```

```
> db.stuff.insert({thing:"apple"})
WriteResult({
  "nInserted" : 0,
```

```

    "writeError" : {
      "code" : 11000,
      "errmsg" : "insertDocument :: caused by :: 11000 E11000 duplicate key error index: test.stuff.$thing_1 dup
key: { : \"apple\" }"
    }
  }
})

```

Option sur la création d'index : sparse

```

> db.produits.insert({item:"polo shirt",size:"medium"})
WriteResult({ "nInserted" : 1 })
> db.produits.insert({item:"jeans",size:"32x32"})
WriteResult({ "nInserted" : 1 })
> db.produits.insert({item:"iphone"})
WriteResult({ "nInserted" : 1 })
> db.produits.insert({item:"DVI-to-VGA cable"})
WriteResult({ "nInserted" : 1 })
> db.produits.find()
{ "_id" : ObjectId("53733414a6d5dcf0dfba69a3"), "item" : "polo shirt", "size" : "medium" }
{ "_id" : ObjectId("5373343aa6d5dcf0dfba69a4"), "item" : "jeans", "size" : "32x32" }
{ "_id" : ObjectId("53733461a6d5dcf0dfba69a5"), "item" : "iphone" }
{ "_id" : ObjectId("5373348fa6d5dcf0dfba69a6"), "item" : "DVI-to-VGA cable" }

```

Problème si on veut poser un index unique sur « size »

```

> db.produits.ensureIndex({size:1},{unique:true})
{
  "createdCollectionAutomatically" : false,
  "numIndexesBefore" : 1,
  "ok" : 0,
  "errmsg" : "E11000 duplicate key error index: test.produits.$size_1 dup key: { : null }",
  "code" : 11000
}

```

null car mongo considère que les deux derniers documents ont le champs « size » égal à null.

Solution : sparse

```

> db.produits.ensureIndex({size:1},{unique:true, sparse:true})
{
  "createdCollectionAutomatically" : false,
  "numIndexesBefore" : 1,
  "numIndexesAfter" : 2,
  "ok" : 1
}

> db.produits.find({size:"medium"}).explain("executionStats")

> db.produits.find().sort({size:1}).explain("executionStats")

```

```

"inputStage" : {
  "stage" : "COLLSCAN",
  "filter" : {

```

Pb : car il existe des documents qui n'ont pas le champ « size »

La commande : explain

```
> mongoimport -d test -c foo < foo.json
```

connected to: 127.0.0.1

2014-05-14T10:31:28.506+0200 check 9 10000

2014-05-14T10:31:28.507+0200 imported 10000 objects

```
> db.foo.find()
```

```

{ "_id" : ObjectId("5373294e620bff2007825235"), "a" : 0, "b" : 0, "c" : 0 }
{ "_id" : ObjectId("5373294e620bff2007825236"), "a" : 1, "b" : 1, "c" : 1 }
{ "_id" : ObjectId("5373294e620bff2007825237"), "a" : 2, "b" : 2, "c" : 2 }
....

```

```
> db.foo.ensureIndex({a:1,b:1,c:1})
```

```

{
  "createdCollectionAutomatically" : false,
  "numIndexesBefore" : 1,
  "numIndexesAfter" : 2,
  "ok" : 1
}

```

Effectuons une requête qui n'utilise pas d'index

```
> db.foo.find({c:1})
```

```
{ "_id" : ObjectId("5373294e620bff2007825236"), "a" : 1, "b" : 1, "c" : 1 }
```

```
> db.foo.find({c:1}).explain("executionStats")
```

```

{
  "executionStats" : {
    "executionSuccess" : true,
    "nReturned" : 1,
    "executionTimeMillis" : 14,
    "totalKeysExamined" : 0,
    "totalDocsExamined" : 10000,
    "executionStages" : {
      "stage" : "COLLSCAN",
      "filter" : {
        "c" : {
          "$eq" : 1
        }
      }
    },
    "nReturned" : 1,
    "executionTimeMillisEstimate" : 0,
    "works" : 10002,
    "advanced" : 1,
    "needTime" : 10000,

```

Effectuons une requête qui utilise l'index

```
> db.foo.find({a:1}).explain("executionStats")
```

```
"nReturned" : 1,  
"executionTimeMillis" : 0,  
"totalKeysExamined" : 1,  
"totalDocsExamined" : 1,
```

```
  "alreadyHasObj" : 0,  
  "inputStage" : {  
    "stage" : "IXSCAN",  
    "nReturned" : 1,  
    "executionTimeMillisEstimate" : 0,  
    "works" : 2
```

```
> db.foo.find({a:500},{a:1,b:1,_id:0})
```

```
{ "a" : 500, "b" : 500 }
```

```
> db.foo.find({a:500},{a:1,b:1,_id:0}).explain("executionStats")
```

```
"executionSuccess" : true,  
"nReturned" : 1,  
"executionTimeMillis" : 0,  
"totalKeysExamined" : 1,  
"totalDocsExamined" : 0,  
"executionStages" : {
```

Maintenant montrons que mongoDB peut utiliser l'index pour la partie tri et pas sur la partie de recherche

```
> db.foo.find({$and:[{c:{$gt:250}},{c:{$lte:500}}]}).sort({a:1}).explain("executionStats")
```

```
"restoreState" : 78,  
"isEOF" : 1,  
"invalidates" : 0,  
"docsExamined" : 10000,  
"alreadyHasObj" : 0,  
"inputStage" : {  
  "stage" : "IXSCAN",  
  "nReturned" : 10000,  
  "executionTimeMillisEstimate" : 30,  
  "works" : 10000,  
  "advanced" : 10000,  
  "needTime" : 0,  
  "needFetch" : 0,
```